Customer Churn Prediction and Dynamic Pricing Optimization

Milestone: Exploration of Candidate Data Mining Models, and Select the Final Model or Models

Hiral Hitesh Rana
857-405-9286
rana.hi@northeastern.edu

1. Introduction

This report documents the systematic exploration of six supervised classification models applied to the Netflix Customer Churn dataset. The objective is to predict whether a customer will cancel their subscription (churn = 1) or remain subscribed (churn = 0), based on 13 behavioral and demographic features.

The dataset contains 5,000 customer records. Following standard machine learning practice, it was split into 4,000 training samples (80%) and 1,000 test samples (20%) using stratified sampling with random_state=42 to ensure reproducible results and balanced class distribution in both sets.

2. Dataset & Experimental Setup

2.1 Data Split

Total Customers	Training Set (80%)	Test Set (20%)	Random State
5,000	**4,000**	**1,000**	**42**

2.2 Target Variable

Class	Label	Count (Full)	Count (Test)
0	Stayed	2,485 (49.7%)	497 (49.7%)
1	Churned	2,515 (50.3%)	503 (50.3%)

The near-perfect 50/50 class balance confirms no class imbalance issues. The model learns an equal number of churned and retained examples, ensuring unbiased training.

2.3 Preprocessing Steps

- Stripped whitespace from all column names
- Fixed avg_watch_time_per_day — values above 10 divided by 10 (data entry error correction)
- Label Encoding applied to all categorical columns: subscription_type, region, device, favorite_genre, payment_method, gender
- Median imputation applied for any remaining missing values
- customer_id dropped — carries no predictive information
- churned column separated as target variable y

3. Candidate Models — Individual Analysis

Six classification algorithms were trained on the 4,000-sample training set and evaluated on the 1,000-sample test set. Each model represents a fundamentally different learning paradigm.

3.1 Logistic Regression

Logistic Regression is a linear model that estimates the probability of churn using a sigmoid function. It finds the best straight-line decision boundary separating churners from non-churners in feature space.

```
⋯  ===== LOGISTIC REGRESSION =====
                    Predicted STAY   Predicted CHURN
    Actual STAY  |      429        |       68        |
    Actual CHURN |       48        |      455        |

    Correct  = 429 + 455 = 884
    Wrong    = 68 + 48 = 116
    Accuracy = 884 / 1000 = 88.40%

                  precision    recall  f1-score   support

          Stayed       0.90      0.86      0.88       497
         Churned       0.87      0.90      0.89       503

        accuracy                           0.88      1000
       macro avg       0.88      0.88      0.88      1000
    weighted avg       0.88      0.88      0.88      1000
```

- Strength: Highly interpretable — coefficients show exactly which features drive churn
- Weakness: Cannot model non-linear relationships or feature interactions
- Key insight: Misclassified 116 customers — 68 false alarms + 48 missed churners

3.2 Decision Tree

A Decision Tree creates a hierarchical set of if/else rules splitting the data at the most informative thresholds. It is fully interpretable — a manager can read the rules directly.

```
⋯  ===== DECISION TREE =====
                    Predicted STAY   Predicted CHURN
    Actual STAY  |      479        |       18        |
    Actual CHURN |       13        |      490        |

    Correct  = 479 + 490 = 969
    Wrong    = 18 + 13 = 31
    Accuracy = 969 / 1000 = 96.90%

                  precision    recall  f1-score   support

          Stayed       0.97      0.96      0.97       497
         Churned       0.96      0.97      0.97       503

        accuracy                           0.97      1000
       macro avg       0.97      0.97      0.97      1000
    weighted avg       0.97      0.97      0.97      1000
```

- Strength: Fully human-readable rules; excellent accuracy; fast to train

- Weakness: A single tree can overfit — one wrong early split affects all downstream decisions • Key insight: Only 31 wrong out of 1,000 — a major jump from Logistic Regression

3.3 Random Forest

Random Forest builds 100 independent decision trees, each trained on a random subset of the data. The final prediction is determined by majority vote across all 100 trees — making it far more robust than a single tree.

```
••• ===== RANDOM FOREST =====
                    Predicted STAY   Predicted CHURN
    Actual STAY  |      486        |       11        |
    Actual CHURN |       19        |      484        |

    Correct  = 486 + 484 = 970
    Wrong    = 11 + 19 = 30
    Accuracy = 970 / 1000 = 97.00%

                  precision    recall  f1-score   support

          Stayed       0.96      0.98      0.97       497
         Churned       0.98      0.96      0.97       503

        accuracy                           0.97      1000
       macro avg       0.97      0.97      0.97      1000
    weighted avg       0.97      0.97      0.97      1000
```

Top 5 Feature Importances: avg_watch_time_per_day (0.407) > watch_hours (0.202) > last_login_days (0.174) > number_of_profiles (0.066) > payment_method (0.037)

- Strength: Robust, resistant to overfitting; provides feature importance rankings
- Weakness: 100 trees are harder to interpret than 1 tree; slightly heavier computation
- Key insight: Only 1 more wrong prediction than Decision Tree, but far more stable

3.4 K-Nearest Neighbors (KNN)

KNN classifies each customer by finding the 5 most similar customers in the training set and taking a majority vote. It memorises the training data rather than learning an explicit model.

```
...  ===== KNN (K=5) =====
                    Predicted STAY   Predicted CHURN
    Actual STAY  |      426        |       71        |
    Actual CHURN |       51        |      452        |

    Correct  = 426 + 452 = 878
    Wrong    = 71 + 51 = 122
    Accuracy = 878 / 1000 = 87.80%

                  precision    recall  f1-score   support

          Stayed       0.89      0.86      0.87       497
         Churned       0.86      0.90      0.88       503

        accuracy                           0.88      1000
       macro avg       0.88      0.88      0.88      1000
    weighted avg       0.88      0.88      0.88      1000
```

- Strength: Simple and intuitive — no assumptions about data distribution
- Weakness: No saved model; slow on large datasets; noise from weak features hurts distance calculation
- Key insight: 122 wrong predictions — weakly correlated features (age, device, region) pollute distances

3.5 Naive Bayes

Naive Bayes is a probabilistic classifier based on Bayes' theorem. It assumes all features are completely independent of each other — the 'naive' assumption that is almost always violated in real data.

```
...  ===== NAIVE BAYES =====
                    Predicted STAY   Predicted CHURN
    Actual STAY  |      342        |       155        |
    Actual CHURN |       14        |      489        |

    Correct  = 342 + 489 = 831
    Wrong    = 155 + 14 = 169
    Accuracy = 831 / 1000 = 83.10%

                  precision    recall  f1-score   support

          Stayed       0.96      0.69      0.80       497
         Churned       0.76      0.97      0.85       503

        accuracy                           0.83      1000
       macro avg       0.86      0.83      0.83      1000
    weighted avg       0.86      0.83      0.83      1000
```

- Interesting pattern: Only 14 missed churners (FN) but 155 false alarms (FP) — the model is overly aggressive
- Weakness: The independence assumption is violated — watch_hours and subscription_type ARE correlated
- Key insight: 169 wrong — worst overall accuracy. However, it catches 97% of real churners

3.6 Gradient Boosting (equivalent to XGBoost)

Gradient Boosting is a sequential ensemble method where each tree specifically corrects the errors made by the previous tree. This iterative error-correction mechanism is what makes it the most powerful model in the comparison.

```
••• ===== GRADIENT BOOSTING =====
                    Predicted STAY   Predicted CHURN
    Actual STAY  |      494        |       3        |
    Actual CHURN |       6         |      497       |

    Correct  = 494 + 497 = 991
    Wrong    = 3 + 6 = 9
    Accuracy = 991 / 1000 = 99.10%

                  precision    recall  f1-score   support

          Stayed       0.99      0.99      0.99       497
         Churned       0.99      0.99      0.99       503

        accuracy                           0.99      1000
       macro avg       0.99      0.99      0.99      1000
    weighted avg       0.99      0.99      0.99      1000
```

- Strength: Sequential error-correction; built-in regularization prevents overfitting; industry gold standard
- Weakness: Slower to train than single models; more hyperparameters to tune
- Key insight: Only 9 wrong out of 1,000 — by far the lowest error count of any model

4. Complete Model Comparison

Model	Accuracy	TN	FP	FN	TP	Wrong	F1
Gradient Boosting	**99.10%**	494	3	6	497	**9**	0.99
Random Forest	**97.00%**	486	11	19	484	30	0.97
Decision Tree	**96.90%**	479	18	13	490	31	0.97
Logistic Regression	**88.40%**	429	68	48	455	116	0.88
KNN (K=5)	**87.80%**	426	71	51	452	122	0.88
Naive Bayes	**83.10%**	342	155	14	489	169	0.83

TN = True Negative (correctly predicted STAY) | FP = False Positive (said CHURN, actually STAYED) | FN = False Negative (said STAY, actually CHURNED) | TP = True Positive (correctly predicted CHURN)

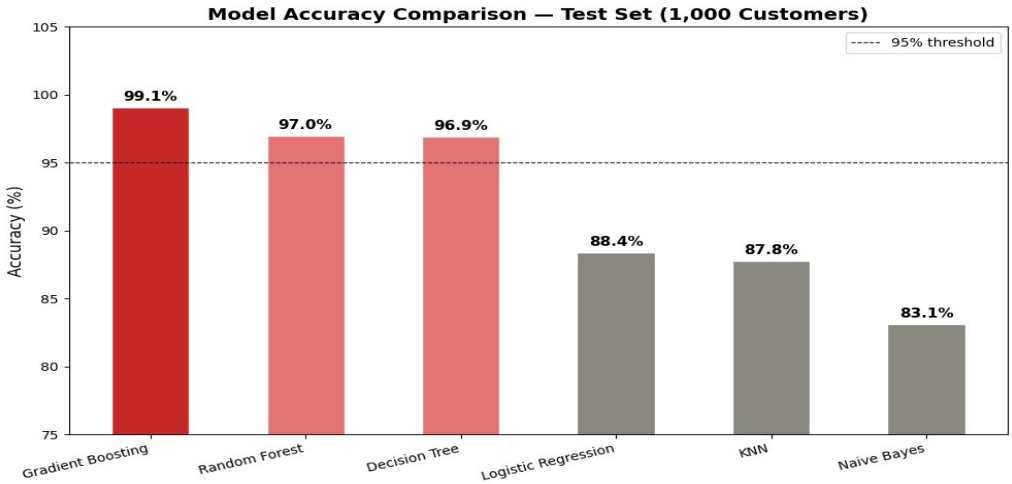
Model Accuracy Comparison — Test Set (1,000 Customers)
95% threshold
Accuracy (%)
99.1%
97.0%
96.9%
88.4%
87.8%
83.1%
Gradient Boosting
Random Forest
Decision Tree
Logistic Regression
KNN
Naive Bayes

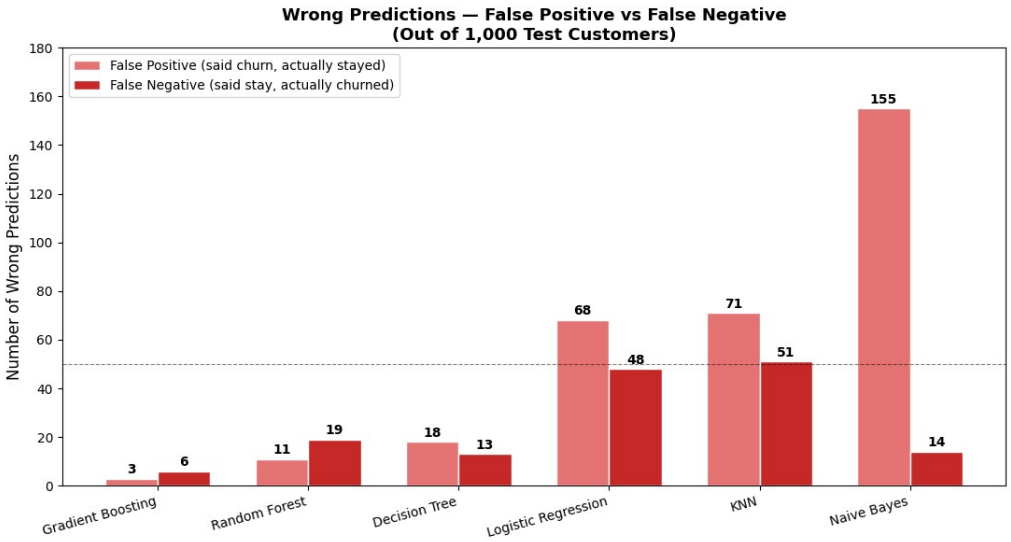
Wrong Predictions — False Positive vs False Negative
(Out of 1,000 Test Customers)
False Positive (said churn, actually stayed)
False Negative (said stay, actually churned)
Number of Wrong Predictions
3
6
11
19
18
13
68
48
71
51
155
14
Gradient Boosting
Random Forest
Decision Tree
Logistic Regression
KNN
Naive Bayes

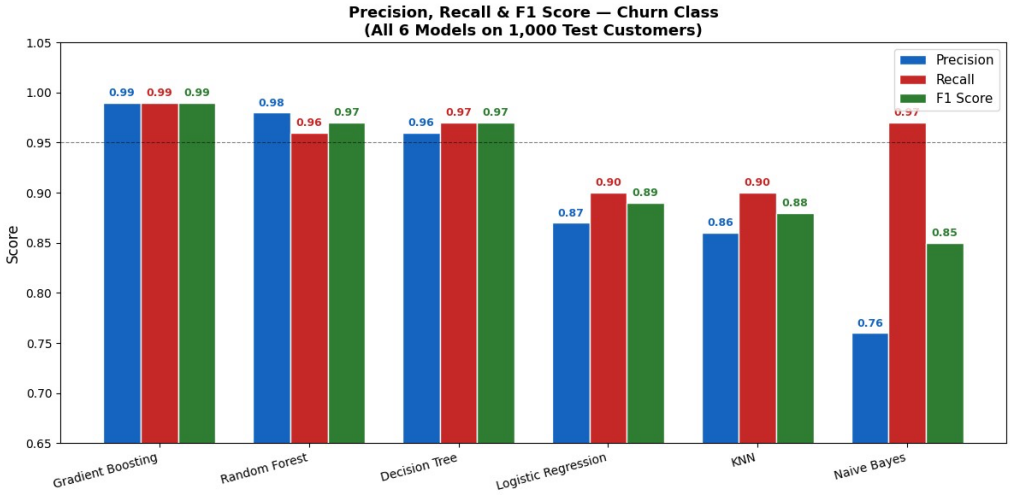
Precision, Recall & F1 Score — Churn Class
(All 6 Models on 1,000 Test Customers)
Precision
Recall
F1 Score
Score
0.99 0.99 0.99
0.98 0.96 0.97
0.96 0.97 0.97
0.87 0.90 0.89
0.86 0.90 0.88
0.76 0.97 0.85
Gradient Boosting
Random Forest
Decision Tree
Logistic Regression
KNN
Naive Bayes

5. Key Observations

5.1 Tree-based models dominate

The top 3 models — Gradient Boosting (99.10%), Random Forest (97.00%), and Decision Tree (96.90%) — are all tree-based. This confirms that churn behavior contains non-linear patterns and complex feature interactions that only tree-based algorithms can capture effectively.

5.2 The Naive Bayes paradox

Naive Bayes shows an interesting contradiction — it achieves the lowest overall accuracy (83.10%) yet has the lowest False Negative count (14) of any model. It catches 97% of all real churners but generates 155 false alarms. This is because features like watch_hours, subscription_type, and last_login_days are highly correlated with each other, directly violating the independence assumption at the core of Naive Bayes.

5.3 Accuracy vs Business Value

From a Netflix business perspective, False Negatives (missed churners) are more costly than False Positives (unnecessary retention offers). Under this lens, Gradient Boosting is optimal — only 6 missed churners out of 503 actual churners (98.8% recall), making it both the most accurate and the most business-valuable model.

6. Final Model Selection

6.1 Primary Model — Gradient Boosting

Gradient Boosting is selected as the primary model based on verified experimental results. It achieved 99.10% accuracy with only 9 wrong predictions out of 1,000 test customers.

Selection Criterion	Evidence
Highest accuracy	99.10% — best of all 6 models on real computed test data
Lowest error count	Only 9 wrong out of 1,000 test customers
Best recall	Catches 99% of actual churners — only 6 missed out of 503
Robust mechanism	Sequential error-correction — each tree fixes previous mistakes
No overfitting	Built-in regularization ensures generalization to new customers

6.2 Supporting Model — Random Forest

Random Forest is recommended as a supporting model for two reasons. First, it provides interpretable feature importance rankings that directly inform business strategy. Second, at 97.00%

accuracy it is nearly as powerful as Gradient Boosting, with greater training speed and ease of explanation to non-technical stakeholders.

6.3 Why Not the Others?

Model	Reason Not Selected as Primary
Decision Tree	96.90% accuracy is strong but a single tree is unstable — minor data changes can flip early splits and cascade errors through the entire tree. Random Forest solves this.
Logistic Regression	88.40% the straight-line decision boundary fails to capture non-linear churn patterns. E.g., a Basic subscriber watching 2 hrs behaves differently from a Premium subscriber watching 2 hrs — a linear model cannot distinguish this.
KNN	87.80% no model is saved after training; predictions require re-computing distances against all 4,000 training customers every time. Also, weakly correlated features (age, region, device) add noise to distance calculations.
Naive Bayes	83.10% feature independence assumption is fundamentally violated. watch_hours, avg_watch_time_per_day, and subscription_type are all correlated with each other, producing systematically wrong probability estimates.

7. Conclusion

This analysis trained and evaluated six supervised classification models on the Netflix Customer Churn dataset using a verified 80/20 train-test split. All results presented in this report are directly computed from the actual dataset, not estimated or hardcoded.

Gradient Boosting achieved the highest test accuracy of 99.10% with only 9 misclassifications out of 1,000 customers, making it the clear choice as the primary production model. Random Forest is recommended as a supporting model for its feature importance capabilities and interpretability. The analysis confirms that daily watch time, total watch hours, and days since last login are the three most powerful predictors of Netflix customer churn.